

CSE 559A: Computer Vision



[credit: danjodon.deviantart.com]

Fall 2017: T-R: 11:30-1pm @ Lopata 101

Instructor: Ayan Chakrabarti (ayan@wustl.edu).

Staff: Abby Stylianou (abby@wustl.edu), Jarett Gross (jarett@wustl.edu)

<http://www.cse.wustl.edu/~ayan/courses/cse559a/>

Sep 5, 2017

OFFICE HOURS

Jarett Gross	Mon	5:40pm-6:30pm	@ TBD
Ayan Chakrabarti	Wed	9:30am-10:30am	@ Jolley 205
Abby Stylianou*	Fri	10:00am-11:00am	@ TBD

Mon/Fri locations will be decided in a day or two.

* Some of the Friday slots will be allocated as recitation sections (one for each problem set).
Dates will be posted in advance.

CONVENTION

RECAP

- An image X is an *array** of intensities.
- $X[n]$ or $X[n_x, n_y]$ refers to intensities for a particular pixel at location n or $[n_x, n_y]$.
 - Single index denotes $n = [n_x, n_y]^T$ is a vector of two integers.
- Each $X[n]$ is a scalar for a grayscale image, or a 3-vector for an RGB color image.
(Unless otherwise specified, vector implies column vector)

$$Y[n] = \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix} X[n]$$

$Y[n]$ $X[n]$


*Clarification: numpy convention is H x W x C: (vertical, horizontal, channels) or H x W.

Do not think of single-channel images themselves as matrices !

It makes no sense to "matrix multiply" a 80x60 pixel image with a 60x20 pixel image.

CONVENTION: LINEAR OPERATIONS

- But sometimes, we want to interpret operations as linear on all intensities / intensity vectors in an image.
- Stack all pixel locations, in some pre-determined order, as rows. Represent X as:
 - $(HW) \times 3$ matrix: color images
 - $(HW) \times 1$ vector: grayscale images.



$$Y[n] = C X[n] \Rightarrow Y = X C^T$$

```
# Begin with X as (H,W,3) array
Xflt = np.reshape(X, (-1,3))    # Flatten X to a (H*W, 3) matrix
Yflt = np.matmul(Xflt,C.T)      # Post-multiply by C
Y = np.reshape(Yflt,X.shape)    # Turn Y back to an image array
```

CONVOLUTION

Notation: $Y = X * k$


$$Y[n] = \sum_{n'} k[n'] X[n - n']$$

$$Y[n_x, n_y] = \sum_{n'_x} \sum_{n'_y} k[n'_x, n'_y] X[(n_x - n'_x), (n_y - n'_y)]$$


- Double summation over the support / size of the kernel k
- We assume $k[n] \in \mathbb{R}$ is scalar valued.
 - If $X[n]$ is scalar, so is $Y[n]$.
 - If X is a color image, each channel convolved with k independently.

To go from m to n channels in a "conv layer": $k[n] \in \mathbb{R}^{n \times m}$ is matrix valued, and $k[n'] X[n - n']$ is a matrix-vector product.


CONVOLUTION



$X[n]$




$Y[n]$



This assumes a 0 centered kernel

$$Y[n] = \sum_{n'} k[n'] X[n - n']$$


CONVOLUTION



We pass 2D arrays to the convolve functions, and get a 2D array out. Let's assume top left index is $(0,0)$ for all.




Let W_x , W_k and W_y denote the widths of X , k , and Y ; and H_x , H_k and H_y the heights.

The 2D convolution function in most libraries provide 3 options: Valid, Full, and Same.



$$Y[n] = \sum_{n'} k[n'] \quad X[n - n']$$

CONVOLUTION






$$W_y = W_x - W_k + 1$$

$$H_y = H_x - H_k + 1$$

Valid: Subet of values of $Y[n]$ for which EVERY $X[n - n']$ is defined.

$$Y[n] = \sum_{n'} k[n'] \quad X[n - n']$$


CONVOLUTION



Full: Subet of values of $Y[n]$ for which ANY $X[n - n']$ is defined.

$$Y[n] = \sum_{n'} k[n'] X[n - n']$$

CONVOLUTION






Padding

What do we use for the missing values of $X[n]$?

- Zero (Often Default)
- Some other constant
- Reflect / Symmetric (across boundary)
- Circular (wrap around)
- Replicate

...

CONVOLUTION



Same: Center Crop of Full output, that is same size as X .

For odd sized kernels, corresponds to treating center of kernel as $(0,0)$.

Same does what we "expect", but you should understand the padding and cropping involved. When kernel size isn't odd, which crop is taken often depends on the library.

CONVOLUTION: PROPERTIES

Let $X *_{\text{full}} k$, $X *_{\text{val}} k$, and $X *_{\text{same}} k$ denote full, valid, and same convolution (with zero padding for full and same)

- **Linear / Distributive:** For scalars α, β ;

- If $Y = X * k$, then: $X * (\alpha k) = (\alpha X) * k = \alpha Y$
- If $Y_1 = X * k_1$ and $Y_2 = X * k_2$, (k_1, k_2 same size): $X * (\alpha k_1 + \beta k_2) = \alpha Y_1 + \beta Y_2$
- If $Y_1 = X_1 * k$ and $Y_2 = X_2 * k$, (X_1, X_2 same size): $(\alpha X_1 + \beta X_2) * k = \alpha Y_1 + \beta Y_2$

- **Associative**

- $(X *_{\text{full}} k_1) *_{\text{full}} k_2 = X *_{\text{full}} (k_1 *_{\text{full}} k_2)$
- $(X *_{\text{val}} k_1) *_{\text{val}} k_2 = X *_{\text{val}} (k_1 *_{\text{full}} k_2)$
- $(X *_{\text{same}} k_1) *_{\text{same}} k_2 \neq X *_{\text{same}} (k_1 *_{\text{full}} k_2)$

- **Commutative:** $k_1 *_{\text{full}} k_2 = k_2 *_{\text{full}} k_1$

- $(X *_{\text{full}} k_1) *_{\text{full}} k_2 = (X *_{\text{full}} k_2) *_{\text{full}} k_1$
- $(X *_{\text{val}} k_1) *_{\text{val}} k_2 = (X *_{\text{val}} k_2) *_{\text{val}} k_1$
- $(X *_{\text{same}} k_1) *_{\text{same}} k_2 \neq (X *_{\text{same}} k_2) *_{\text{same}} k_1$

CONVOLUTION

$$X[n]$$



$$Y[n]$$



0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0
0	0	0	0	0

$$k[n]$$

(Same with zero padding)

CONVOLUTION

$X[n]$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$Y[n]$

Kernel = Impulse Response

$k[n]$

(Same with zero padding)

CONVOLUTION

$$X[n]$$



$$Y[n]$$



0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	1

$$k[n]$$


(Same with zero padding)

CONVOLUTION

$$X[n]$$



$$Y[n]$$



$$k[n] = 1 / 25$$

(Same with zero padding)

CONVOLUTION

$X[n]$



$Y[n]$



$\sigma = 1$

Gaussian Kernels

$$G_{\sigma}[n_x, n_y] \propto \exp\left(-\frac{n_x^2 + n_y^2}{2\sigma^2}\right)$$

$$\sum_{n_x, n_y} G_{\sigma}[n_x, n_y] = 1$$

$$n_x, n_y = [-S, -(S-1), \dots, -1, 0, 1, \dots, (S-1), S]$$

CONVOLUTION

$X[n]$



$Y[n]$



$\sigma = 2$

Gaussian Kernels

$$G_{\sigma}[n_x, n_y] \propto \exp\left(-\frac{n_x^2 + n_y^2}{2\sigma^2}\right)$$

$$\sum_{n_x, n_y} G_{\sigma}[n_x, n_y] = 1$$

$$n_x, n_y = [-S, -(S-1), \dots, -1, 0, 1, \dots, (S-1), S]$$

CONVOLUTION

$X[n]$



$Y[n]$



$\sigma = 3$

Gaussian Kernels

$$G_{\sigma}[n_x, n_y] \propto \exp\left(-\frac{n_x^2 + n_y^2}{2\sigma^2}\right)$$

$$\sum_{n_x, n_y} G_{\sigma}[n_x, n_y] = 1$$

$$n_x, n_y = [-S, -(S-1), \dots, -1, 0, 1, \dots, (S-1), S]$$

CONVOLUTION

$X[n]$



$Y[n]$



$\sigma = 4$

Gaussian Kernels

$$G_{\sigma}[n_x, n_y] \propto \exp\left(-\frac{n_x^2 + n_y^2}{2\sigma^2}\right)$$

$$\sum_{n_x, n_y} G_{\sigma}[n_x, n_y] = 1$$

$$n_x, n_y = [-S, -(S-1), \dots, -1, 0, 1, \dots, (S-1), S]$$

CONVOLUTION

$X[n]$



$Y[n]$



Gaussian Kernels

$$G_{\sigma}[n_x, n_y] \propto \exp\left(-\frac{n_x^2 + n_y^2}{2\sigma^2}\right)$$

$$\sum_{n_x, n_y} G_{\sigma}[n_x, n_y] = 1$$

$$n_x, n_y = [-S, -(S-1), \dots, -1, 0, 1, \dots, (S-1), S]$$

CONVOLUTION

$X[n]$



$Y[n]$



$\sigma = 4 \quad \alpha = 1$

Unsharp Masking

$$Y = (1 + \alpha)X - \alpha(X * G_{\sigma}) = X * ((1 + \alpha)\delta - \alpha G_{\sigma})$$

CONVOLUTION

$X[n]$



$Y[n]$



$\sigma = 2 \quad \alpha = 1$

Unsharp Masking

$$Y = (1 + \alpha)X - \alpha(X * G_{\sigma}) = X * ((1 + \alpha)\delta - \alpha G_{\sigma})$$

CONVOLUTION

$X[n]$



$Y[n]$



$\sigma = 2 \quad \alpha = 5$

Unsharp Masking

$$Y = (1 + \alpha)X - \alpha(X * G_{\sigma}) = X * ((1 + \alpha)\delta - \alpha G_{\sigma})$$

CONVOLUTION

$X[n]$



$Y[n]$







$\sigma = 2 \quad \alpha = 10$

Unsharp Masking

$$Y = (1 + \alpha)X - \alpha(X * G_{\sigma}) = X * ((1 + \alpha)\delta - \alpha G_{\sigma})$$

APPLICATION: EDGE DETECTION

What is an edge ?



Different answers from different people !

Depth boundary / Material Boundary / Object Boundary ?

Edge (not boundary): Location where image intensity is changing rapidly in some direction.

Directional Derivative

APPLICATION: EDGE DETECTION

Finite Difference Approximation

$$\frac{\partial}{\partial n_x} X[n_x, n_y] \propto X[n_x + 1, n_y] - X[n_x - 1, n_y]$$

$$X * [1 \ 0 \ -1]$$

Derivative is a linear spatially invariant operation: Convolution

$$X * \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

Smoothed in y direction
"Sobel" Operator


$$X * \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Y Derivative

APPLICATION: EDGE DETECTION



$$X * \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$



$$X * \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Derivatives have been scaled so that gray (0.5) corresponds to 0. Bright to positive derivative values, dark to negative.

APPLICATION: EDGE DETECTION

Smoothing + Derivative

$$I_x = \partial_x * (G_\sigma * X) = (\partial_x * G_\sigma) * X = G_{x:\sigma} * X$$

$$G_x = \frac{-x}{2\pi\sigma^4} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) \quad G_y = \frac{-y}{2\pi\sigma^4} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

Derivative of Gaussian (DoG) Filters

APPLICATION: EDGE DETECTION

Smoothing + Derivative

$$I_x = \partial_x * (G_\sigma * X) = (\partial_x * G_\sigma) * X = G_{x:\sigma} * X$$

$$G_x = \frac{-x}{2\pi\sigma^4} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) \quad G_y = \frac{-y}{2\pi\sigma^4} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

$$G_\theta = \frac{-(x \cos \theta + y \sin \theta)}{2\pi\sigma^4} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

$$I_\theta = I_x \cos \theta + I_y \sin \theta$$

Just need to convolve twice. Gives us an expression for derivative along every direction.

APPLICATION: EDGE DETECTION

Smoothing + Derivative

$$I_{\theta}[n] = I_x[n] \cos \theta + I_y[n] \sin \theta$$

$$H[n] = \sqrt{I_x^2[n] + I_y^2[n]} = \max_{\theta} I_{\theta}[n]$$

$$\Theta[n] = \text{atan2}(I_y, I_x) = \arg \max_{\theta} I_{\theta}[n]$$


Gives us gradient magnitude and direction.

Often applied even to filters that aren't "steerable" like DoG.


APPLICATION: EDGE DETECTION



I_x



I_y




I_{45°

APPLICATION: EDGE DETECTION



I_x



I_y




H

APPLICATION: EDGE DETECTION



I_x



I_y




$H > \epsilon$

APPLICATION: EDGE DETECTION



I_x



I_y




$H > \epsilon$

APPLICATION: EDGE DETECTION

Extensions

- Non-maxima Supression: Keep an edge pixel only if its magnitude is higher than its neighbors *along the direction of the derivative*.



$H[n]$


Declare edge if a above threshold and :

- $a > b$ and $a > c$ if $\theta = 0$
- $a > f$ and $a > j$ if $\theta = 90$
- $a > e$ and $a > k$ if $\theta = 45$
-

- Canny: Keep a lower magnitude edge pixel if it has a higher edge magnitude neighbor.
Two thresholds (hysteresis)
- Second derivative filters.

See Szeliski Section 4.2

LINES




Missed detections
Clutter
Occlusions

Pool them together to detect
scene structure: E.g., Lines

Edges are isolated per-pixel labels


LINES



Hough Transform

- Consider ALL possible lines (on a 2D plane)

LINES



Hough Transform

- Consider ALL possible lines (on a 2D plane)
- This is a two dimensional search space, could parameterize it in different ways.


$$r = x \cos \theta + y \sin \theta$$

$$\theta \in [-\pi/2, \pi/2]$$


$$r \in [-r_{\max}, r_{\max}]$$

LINES

$$r = x \cos \theta + y \sin \theta$$



Hough Transform





Discretize this space into a bunch of buckets.

LINES

$$r = x \cos \theta + y \sin \theta$$

Hough Transform




Each edge pixel casts a vote for all lines it could belong to.

Compute by plugging in x, y into equation: get a value of r for each θ (get a sinusoid)

LINES

$$r = x \cos \theta + y \sin \theta$$

Hough Transform



Dominant Lines

Each edge pixel casts a vote for all lines it could belong to.

Compute by plugging in x,y into equation: get a value of r for each θ (get a sinusoid)

Do this for all pixels and see which 'bins' get the most votes.

Variants


- Each edge pixel only casts one vote based on angle. (with or without sign)
- Vote weighted by magnitude of gradient.
- Exclusive vote (select dominant line, remove vote from its pixels for other lines)
- Use same idea for line segments, circles, ellipses ...

OTHER NEIGHBORHOOD OPERATIONS

Median Filter / Order Statistics

$$Y[n] = \text{Median}\{X[n - n']\}_{N[n']=1}$$

- Neighborhood function $N[n'] \in \{0, 1\}$
- Often better at removing outliers than convolution.



Source: Wikipedia

- Other ops: $Y[n] = \max / \min\{X[n - n']\}_{N[n']>0}$

OTHER NEIGHBORHOOD OPERATIONS

Morphological Operations

- Conducted on binary images ($X[n] \in \{0, 1\}$)
- Erosion: $Y[n] = \text{AND } \{X[n - n']\}_{N[n']=1}$ (1 if all neighbors 1)
- Dilation: $Y[n] = \text{OR } \{X[n - n']\}_{N[n']=1}$ (1 if any neighbor 1)
- Opening: Erosion followed by Dilation
- Closing: Dilation followed by Erosion

See Szeliski Sec 3.3.2




Figure 3.21 Binary image morphology: (a) original image; (b) dilation; (c) erosion; (d) majority; (e) opening; (f) closing. The structuring element for all examples is a 5×5 square. The effects of majority are a subtle rounding of sharp corners. Opening fails to eliminate the dot, since it is not wide enough.

NEXT TIME

- Bilateral Filtering
- Fourier Transforms
- Making convolutions efficient
- Sampling and Scale
- Image Representations