

# CSE 559A: Computer Vision



[credit: danjodon.deviantart.com]

Fall 2017: T-R: 11:30-1pm @ Lopata 101

Instructor: Ayan Chakrabarti (ayan@wustl.edu).

Staff: Abby Stylianou (abby@wustl.edu), Jarett Gross (jarett@wustl.edu)

<http://www.cse.wustl.edu/~ayan/courses/cse559a/>

Sep 12, 2017


# ADMINISTRIVIA

---

- Homework becomes available at 5pm today.
- Posted on blackboard. Submission also on blackboard.
- If you are sitting in and would like a copy, e-mail me. (won't be graded)
- Lots of support code (you fill in specific functions). Read the support code !
- $\text{\textit{L}A\text{\textit{T}}E\text{\textit{X}}}$  template for report is now in the resources section, along with tutorials/links.
- Please finish information section (discussions, online sources, number of hours spent)


# SCALE & ALIASING

---



$W \times H$

"Resize" Images




$(W/2) \times (H/2)$




# SCALE & ALIASING

---




$W \times H$



$(W/2) \times (H/2)$

"Resize" Images




[Source: Wikipedia]


"Aliasing"

# SCALE & ALIASING


---




Fourier Transform  
→



If you write it out,  
you see the higher  
freq. components  
get folded into  
lower freq.




Fourier Transform  
→





Remember, in the two cases  $F[u,v]$  is defined  
with respect to different width and height  
 $W_x$  and  $H_x$ , and for different ranges of  $(u,v)$ .

# SCALE & ALIASING


---



Fourier Transform  
→



Fourier Transform  
→



Make sure there are  
no high frequencies  
before sub-sampling !

Low-pass filter, i.e.,  
Smooth Image before  
sub-sampling.

# SCALE & ALIASING

---



Without Smoothing




With Smoothing


Sometimes the camera itself makes aliased measurements: if spatial sensitivity is low at edges of pixel.

# SCALE & ALIASING

---



W x H



$(W/2) \times (H/2)$

"Resize" Images


- Need to hallucinate missing information.
- Lots of research (super-resolution).
- Simplest Approach: Nearest neighbor

$$Y[n] = X[\text{round}(n/2)]$$




# SCALE & ALIASING

---




$W \times H$



$(W/2) \times (H/2)$

"Resize" Images


- Need to hallucinate missing information.
- Lots of research (super-resolution).
- Simplest Approach: Nearest neighbor
- Simple Approach: (Bi) Linear Interpolation



For up-sampling by 2 in 1-D, missing values are just the average of the left and right present values.

# SCALE & ALIASING

---



1/4	1/2	1/4
1/2	1	1/2
1/4	1/2	1/4

Can achieve this by filling with zeros, and convolution with a 3x3 kernel.

# EFFICIENT COMPUTATION

---

- Convolution, in the most general case, takes  $O(n_x n_k)$  time.
  - $n_x = W_x H_x, n_k = W_k H_k$ .
- Convolution in the frequency domain:
  - FFT, point-wise multiply, Inverse FFT
  - FFT/IFFT complexity is  $O(n_x \log n_x)$  (Most efficient for power of 2 image size)
  - May be worth it for large kernels
  - Or same image convolved with many different kernels

# EFFICIENT COMPUTATION

---

## Separable Kernels

$$G[n_x, n_y] \propto \exp\left(-\frac{n_x^2 + n_y^2}{2\sigma^2}\right) = G_x[n_x]G_y[n_y]$$

- $x$ - and  $y$ - derivatives of Gaussian also separable.
- Realize that  $k[n_x, n_y] = k_x[n_x]_y k[n_y] = k_x *_{\text{full}} k_y$ .

This is by interpreting  $k_x$  and  $k_y$  as having size  $W_k \times 1$  and  $1 \times H_k$ .

- So  $X * k = X * (k_x * k_y) = (X * k_x) * k_y$  This takes  $W_k + H_k$  operations instead of  $W_k H_k$ .
- Often if a kernel itself isn't separable, it can be sometimes expressed as a sum of separable kernels.
- E.g., Unsharp Mask:  $(1 + \alpha)\delta - \alpha G_\sigma$  (don't combine!)
- Could also try to do this automatically using SVD.

# EFFICIENT COMPUTATION

---

## Recursive Computation

$$\begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 0 & 0 \\ 1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$$
$$\begin{bmatrix} -1 & -1 & 1 & 1 \\ -1 & -1 & 1 & 1 \\ 1 & 1 & -1 & -1 \\ 1 & 1 & -1 & -1 \end{bmatrix} = \begin{bmatrix} -1 & 0 & 1 \\ 0 & 0 & 0 \\ 1 & 0 & -1 \end{bmatrix} * \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$$


- Sometimes can decompose into convolution with sparse kernels.
- Many implementations of `convolve2d` won't make use of sparsity.
  - But you can write your own.

# EFFICIENT COMPUTATION

---

## Smooth and Sub-sample

- Don't smooth and subsample !
- For sub-sampling by two, you're computing 4x as many smooth filter outputs as you need to.



- Similarly, using zero-filling + convolution for upsampling is inefficient.

# EFFICIENT COMPUTATION

---

## numpy Specifics

1. In general, prefer algorithms that have lower total number of multiplies / adds.
  2. Try to use `scipy.signal.convolve2d` (subject to rule 1). It is optimized for cache reads, parallel execution, etc.  
(I import it often as: `from scipy.signal import convolve2d as conv2`)
  3. Similarly, avoid `for` loops and use element-wise operations on large arrays, matrix multiplies (`np.matmul` / `np.dot`), etc.
- Some of these things are faster in python because a single large operation runs natively instead of returning to the compiler. But they're also faster because these operations are often 'atomics' in lower-level libraries too (BLAS), and have been highly optimized for modern hardware.
  - Thinking about designing your algorithm in terms of these atomic operations is useful beyond python.
  - Some points in problem sets allocated for efficient code.


# MULTI-SCALE REPRESENTATIONS

---

Fourier Transform useful for:

Diagonalizing Convolution  $A_k = SD_kS^*$

Concentrating energy (high magnitudes) in fewer coefficients



But is the FT interpretable ? Does it tell us something about the image ?



# MULTI-SCALE REPRESENTATIONS

---

- $F[u, v]$  is intuitively average variation in image at that frequency.
- But averaged across the entire image.
- This isn't useful because images aren't "stationary"
- Different parts of the image, "have different frequencies".

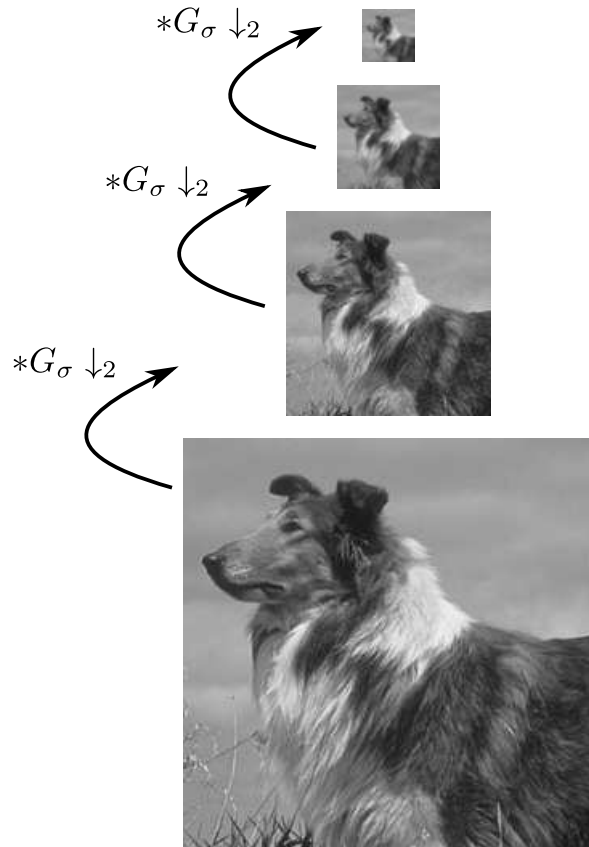


- FT decomposition of different levels (coarse/fine) of variation: but without sense of spatial location.
- Multi-scale representations aim to address this.

# MULTI-SCALE REPRESENTATIONS

---

## Gaussian Pyramid

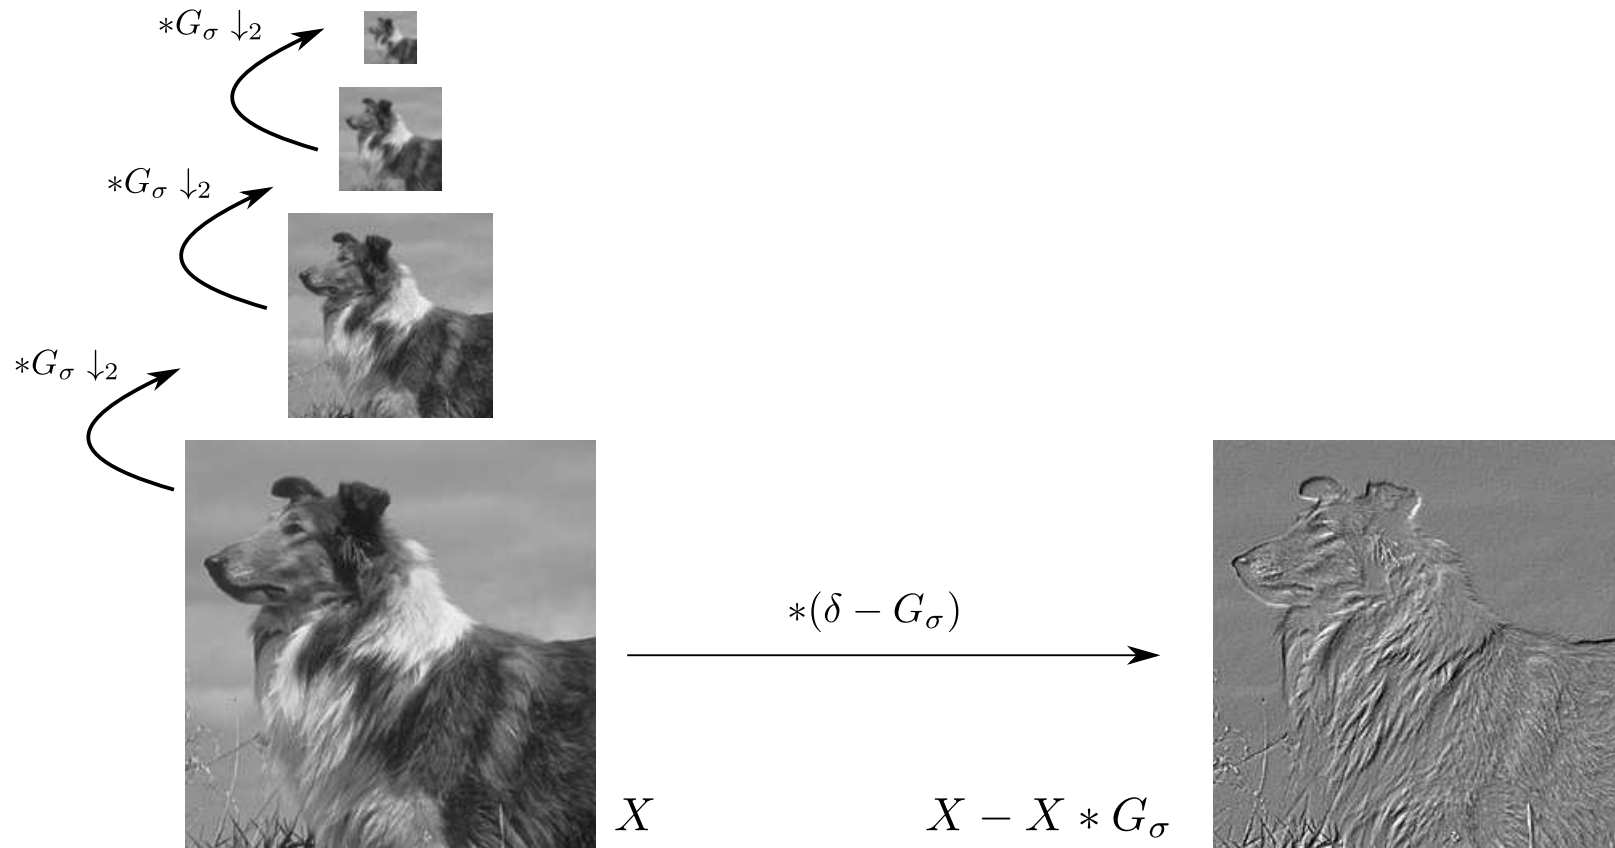


Useful for analyzing image at multiple scales.

E.g., apply the same method (edge detection / CNN) at multiple levels of the pyramid.

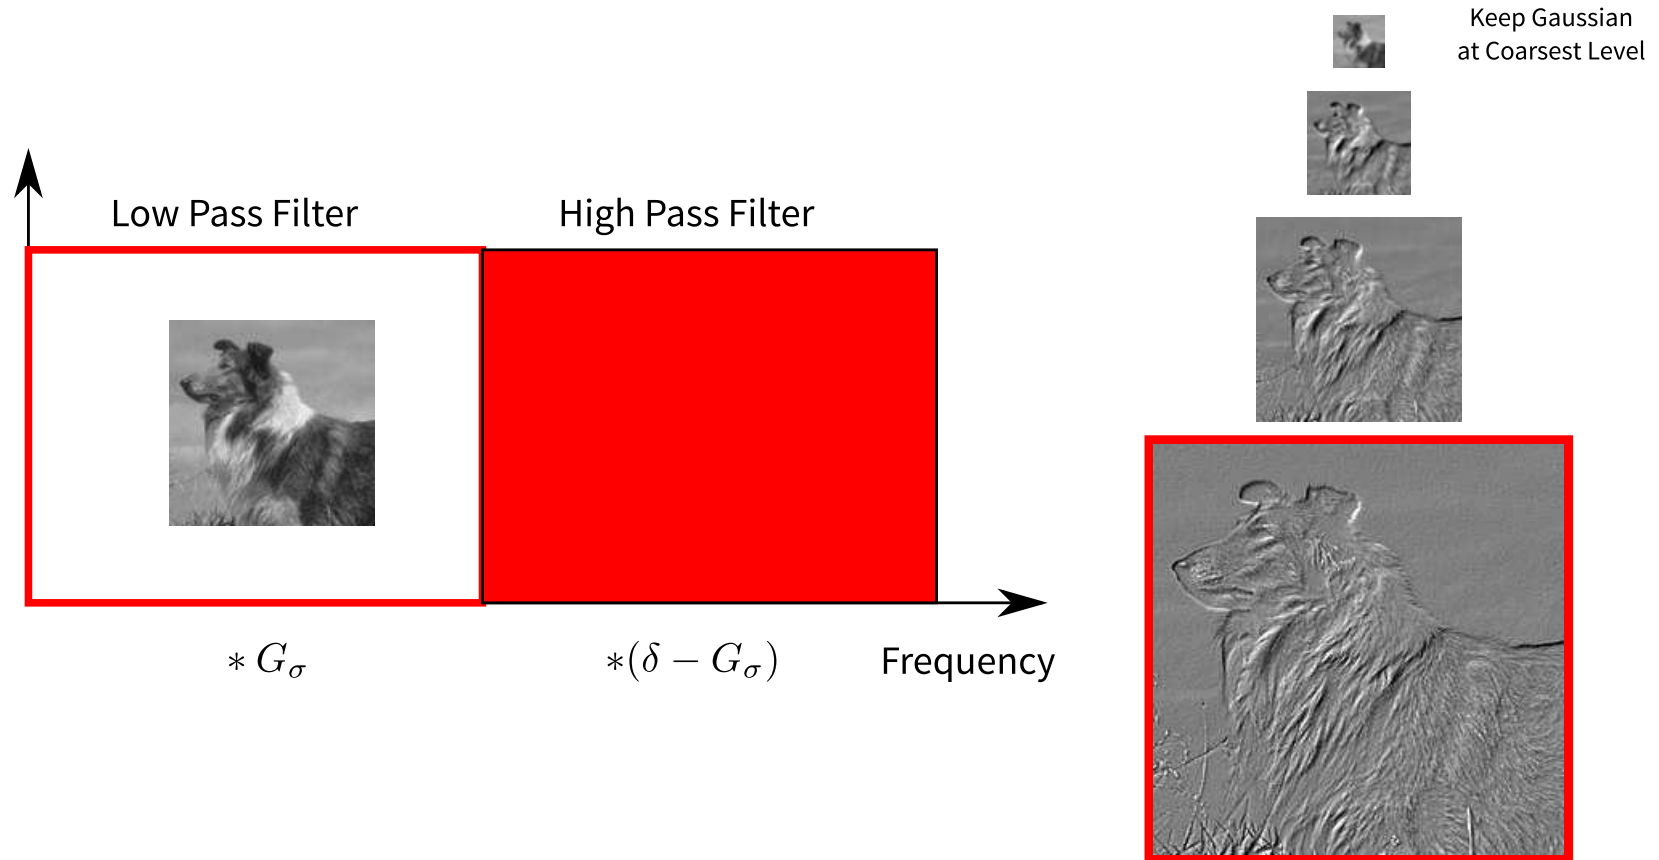
# MULTI-SCALE REPRESENTATIONS

## Laplacian Pyramid



# MULTI-SCALE REPRESENTATIONS

## Laplacian Pyramid


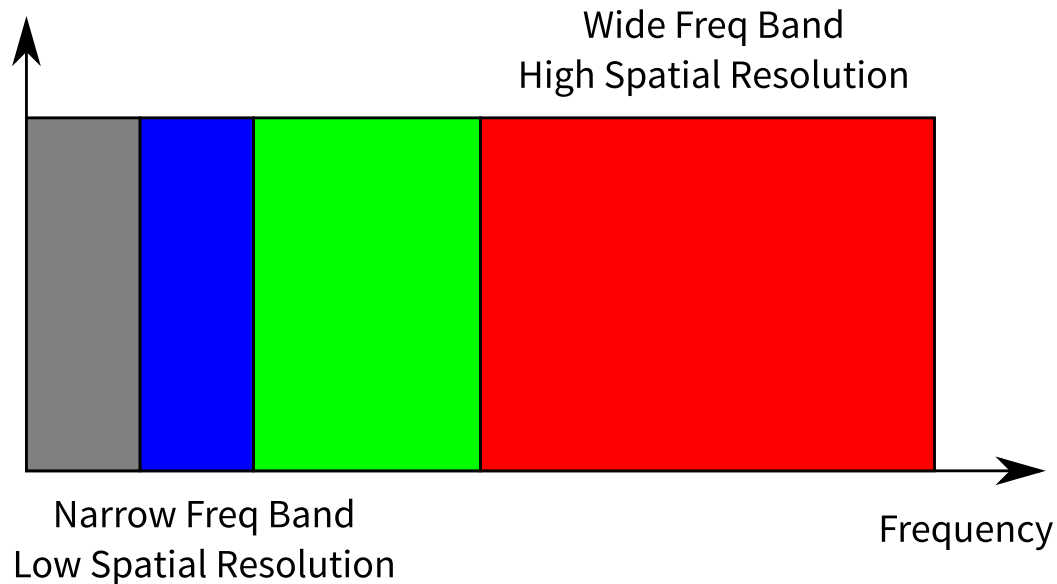


# MULTI-SCALE REPRESENTATIONS

## Laplacian Pyramid

No orientation selectivity: See Steerable Pyramids

<http://www.cns.nyu.edu/~eero/steerpyr/>



# MULTI-SCALE REPRESENTATIONS

---

## Wavelet Pyramid

Gaussian & Laplacian pyramids are good for analysis, but not really a representation.

Not easy / possible to go from pyramid coefficients back to original image (like Fourier Transform).

Consider a 2x2 Pixel Block

a	c
b	d

$$L = \frac{a + b + c + d}{2}$$

$$H_2 = \frac{c + d - a - b}{2}$$

$$H_1 = \frac{b + d - a - c}{2}$$

$$H_3 = \frac{a + d - b - c}{2}$$

$$\begin{bmatrix} L \\ H_1 \\ H_2 \\ H_3 \end{bmatrix} = \frac{1}{2} \begin{bmatrix} 1 & 1 & 1 & 1 \\ -1 & 1 & -1 & 1 \\ -1 & -1 & 1 & 1 \\ 1 & -1 & -1 & 1 \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \end{bmatrix}$$

Unitary Matrix / Co-ordinate transform

# MULTI-SCALE REPRESENTATIONS

---

## Wavelet Pyramid




For every non-overlapping  
2x2 Patch in the Image

a	c
b	d

Can be achieved by  
Convolution + Downsample

Divided by 2  
for Visualization



Still a Co-ordinate Transform !

$$L = \frac{a + b + c + d}{2}$$

$$H_2 = \frac{c + d - a - b}{2}$$

$$H_1 = \frac{b + d - a - c}{2}$$

$$H_3 = \frac{a + d - b - c}{2}$$


# MULTI-SCALE REPRESENTATIONS

## Wavelet Pyramid



$$\begin{bmatrix} L \\ H_1 \\ H_2 \\ H_3 \end{bmatrix} = \frac{1}{2} \begin{bmatrix} 1 & 1 & 1 & 1 \\ -1 & 1 & -1 & 1 \\ -1 & -1 & 1 & 1 \\ 1 & -1 & -1 & 1 \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \end{bmatrix}$$

← Can Invert to get Image



a	c
b	d

$$L = \frac{a + b + c + d}{2}$$

$$H_2 = \frac{c + d - a - b}{2}$$

$$H_1 = \frac{b + d - a - c}{2}$$

$$H_3 = \frac{a + d - b - c}{2}$$




# MULTI-SCALE REPRESENTATIONS

---

## Wavelet Pyramid



Apply  
Recursively



## Wavelet Transform

a	c
b	d

$$L = \frac{a + b + c + d}{2}$$

$$H_1 = \frac{b + d - a - c}{2}$$

$$H_2 = \frac{c + d - a - b}{2}$$

$$H_3 = \frac{a + d - b - c}{2}$$


"Harr"

Others based on  
different filters

# MULTI-SCALE REPRESENTATIONS

---

## Wavelet Pyramid



Applications: Analysis, image modeling & restoration, **compression**

# IMAGE RESTORATION

---

- Recovering "true" image  $X[n]$  from observed image  $Y[n]$   
(can apply to image like objects too, like depth maps)
- Step 1: Assume we know the degradation model, or the mapping from  $X$  to  $Y$

$p(Y|X)$  : Distribution of possible  $Y$ s that we can get from  $X$

Example: Additive White Gaussian Noise

$$Y[n] = X[n] + \epsilon[n], \quad \epsilon[n] \sim \mathcal{N}(0, \sigma^2)$$

$$p(Y[n]|X[n]) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(Y[n] - X[n])^2}{2\sigma^2}\right) \Rightarrow p(Y|X) = \prod_n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(Y[n] - X[n])^2}{2\sigma^2}\right)$$

$$p(Y|X) = (2\pi\sigma^2)^{-N/2} \exp\left(-\frac{\|Y - X\|^2}{2\sigma^2}\right), \quad N = \text{Total no. of pixels.}$$

Example: Blur (+ noise)

$$Y[n] = (X * k)[n] + \epsilon[n], \quad \epsilon[n] \sim \mathcal{N}(0, \sigma^2)$$

For fronto-parallel scenes, defocus and (parallel) motion blur can be modeled as convolution.

# IMAGE RESTORATION

---

## Bayesian View

- Step 2: Define a prior distribution  $p(X)$ , which encodes your *a-priori* knowledge about statistics of natural images.
- Bayes Rule: Given observation likelihood and prior, gives us *Posterior distribution*

$$p(X|Y) = \frac{P(Y|X)P(X)}{\int_{X'} P(Y|X')P(X')dX'} \propto P(Y|X)P(X)$$

- Could estimate  $X$  as the mean / mode of this distribution. Let's focus on the maximum. Called the Maximum A Posteriori (MAP) estimate:

$$\begin{aligned}\hat{X} &= \arg \max_X P(Y|X) = \arg \max_X P(Y|X)P(X) \\ &= \arg \min_X -\log P(Y|X) - \log P(X)\end{aligned}$$

# IMAGE RESTORATION

---

- For denoising:

$$\hat{X} = \arg \min_X \frac{\|Y - X\|^2}{2\sigma^2} - \log P(X)$$

- For deblurring:

$$\hat{X} = \arg \min_X \frac{\|Y - A_k X\|^2}{2\sigma^2} - \log P(X)$$

Instead of Bayesian view, can think of minimizing a data cost + 'regularizer' on  $X$ :

$$\hat{X} = \arg \min_X \frac{\|Y - X\|^2}{2\sigma^2} + R(X)$$

$$\hat{X} = \arg \min_X \frac{\|Y - A_k X\|^2}{2\sigma^2} + R(X)$$

# IMAGE RESTORATION

---

Now let's say we were doing denoising, and our regularizer / prior was pixel-wise.

$$R(X) = \sum_n R_n(X[n])$$

We could find the estimate of each  $X[n]$  independently.

$$\hat{X}[n] = \arg \min_x \frac{(Y[n] - x)^2}{2\sigma^2} + R_n(x), \quad \forall n$$

For example,  $R_n(x) = \frac{(x-0.5)^2}{2\sigma_x^2}$  (- log probability for  $\mathcal{N}(0.5, \sigma_x^2)$ )

$$\begin{aligned} X[n] &= \arg \min_x \frac{(Y[n] - x)^2}{2\sigma^2} + \frac{(0.5 - x)^2}{2\sigma_x^2} \\ &= \frac{Y[n]\sigma_x^2 + 0.5\sigma^2}{\sigma_x^2 + \sigma^2} \end{aligned}$$

(Take derivative of cost function set to 0, check second derivative is positive)

# IMAGE RESTORATION

---

But what happens when the function doesn't decompose pixel-wise ?

Example: Regularizer puts quadratic penalty on gradients.

$$R(X) = \frac{\lambda}{2} \sum_n (G_x * X)[n]^2 + (G_y * X)[n]^2$$

$$\hat{X} = \arg \min_X \frac{1}{2\sigma^2} \|Y - X\|^2 + \frac{\lambda}{2} \left[ \|A_{gx}X\|^2 + \|A_{gy}X\|^2 \right]$$

where  $A_{gx}$  and  $A_{gy}$  are matrix form of convolution with x- and y- derivative filters.

Expanding this out:

$$\hat{X} = \arg \min_X X^T Q X - 2 X^T B + c$$

where:

$$Q = \frac{1}{2\sigma^2} I + \frac{\lambda}{2} (A_{gx}^T A_{gx} + A_{gy}^T A_{gy}), \text{ where } I \text{ is } N \times N \text{ identity matrix.}$$

$$B = Y/\sigma^2, \text{ and } c \text{ is independent of } X.$$

# IMAGE RESTORATION

---

$$\hat{X} = \arg \min_X X^T Q X - 2 X^T B + c$$

where:

$$Q = \frac{1}{2\sigma^2} I + \frac{\lambda}{2} (A_{gx}^T A_{gx} + A_{gy}^T A_{gy}), \text{ where } I \text{ is } N \times N \text{ identity matrix.}$$

$$B = Y/\sigma^2, \text{ and } c \text{ is independent of } X.$$

- This is a vector-equivalent of a quadratic form. . . .
- We can find  $X$  by computing derivative of the cost wrt  $X$ , and setting it to 0.

$$2 \hat{Q}X - 2 B = 0 \Rightarrow \hat{X} = Q^{-1}B$$

- All eigen-values of  $Q$  are strictly above 0. (Why ?)
- "Second" derivative of the cost along any direction of  $X$  is positive. So minima, not maxima.
- $Q$  is called a symmetric positive-definite matrix. The cost function is convex along all directions of  $X$ .

The problem:  $Q$  is  $N \times N$ , where  $N$  is the number of pixels !



# IMAGE RESTORATION

---

Useful Reading:

- [https://en.wikipedia.org/wiki/Positive-definite\\_matrix](https://en.wikipedia.org/wiki/Positive-definite_matrix)

- [https://en.wikipedia.org/wiki/Matrix\\_calculus](https://en.wikipedia.org/wiki/Matrix_calculus)

We'll be using the "denominator layout" convention.