

# CSE 559A: Computer Vision



[credit: danjodon.deviantart.com]

Fall 2017: T-R: 11:30-1pm @ Lopata 101

Instructor: Ayan Chakrabarti (ayan@wustl.edu).

Staff: Abby Stylianou (abby@wustl.edu), Jarett Gross (jarett@wustl.edu)

<http://www.cse.wustl.edu/~ayan/courses/cse559a/>

Sep 14, 2017

# ADMINISTRIVIA

---

- Homework posted (and updated!). Make sure you have `pset1V2.zip`.
- Recitation will be NEXT Friday (9/22).
- Regular office hours tomorrow (in J420).

# CRASH COURSE ON OPTIMIZATION

---

- Let  $x$  be a scalar.
- $f(x; \theta)$  is some function of  $x$ , and some other parameters  $\theta$ .
- $\min_x f(x; \theta)$  is the smallest value that  $f$  can take ...
  - For some fixed values of  $\theta$
  - By searching over all possible values of  $x$
  - Is a function of  $\theta$
  - But not of  $x$




$$f(x; a, b, c) = a(x - b)^2 + c$$

$$\min_x f(x; a, b, c) = a + c$$

# CRASH COURSE ON OPTIMIZATION

- $\arg \min_x f(x; \theta)$  is the value of  $x$  for which  $f$  attains its minimum value.
- Same deal for max and arg max.  $\max f = -(\min(-f))$ .
- How do we find  $x$ ?
- If  $\frac{\partial f(x; \theta)}{\partial x} = 0$  at  $x = x'$ , then  $x'$  is an extremum.
  - i.e., *local* minimum or local maximum.
  - Can find which by checking second derivative.

$$\text{Minimum: } \frac{\partial^2 f(x; \theta)}{\partial x^2} > 0; \quad \text{Maximum: } \frac{\partial^2 f(x; \theta)}{\partial x^2} < 0$$



# CRASH COURSE ON OPTIMIZATION

---

- $f(x; a, b, c) = ax^2 + bx + c$
- Only one minima or maxima at  $-b/2a$
- Can see it also by rewriting as  $a\left(x - \frac{-b}{2a}\right)^2 + c - \frac{b^2}{4a}$
- Minimum if  $a > 0$ , Maximum if  $a < 0$

# CRASH COURSE ON OPTIMIZATION

---

- Minimization over multiple variables

$$\arg \min_{x_1, x_2, x_3} f(x_1, x_2, x_3; \theta)$$

$$\arg \min_x f(x; \theta), \quad x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

- Note that output of  $f$ , which you are minimizing, is still scalar valued (a single number).

# CRASH COURSE ON OPTIMIZATION

---

- Generalization of derivative: gradient

$$\nabla_x f(x; \theta) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \frac{\partial f}{\partial x_3} \end{bmatrix}$$

- Also a vector of the same dimensions as  $x$




$$\frac{\partial f}{\partial (\alpha x_1 + \beta x_2 + \gamma x_3)} = \left\langle \nabla_x f, [\alpha, \beta, \gamma]^T \right\rangle$$

- Derived by chain rule
- Tells us about gradient in any direction.
- $y = Ax \Rightarrow (\nabla_y f) = A (\nabla_x f)$
- If we say  $(\nabla_x f) = 0$  at  $x$ , that means every element of the gradient vector is 0.
- And so, the derivative along all "directions" is 0. Then  $x$  is an extremum of  $f$ .

# CRASH COURSE ON OPTIMIZATION

---

- Identities
  - $\nabla_x x^T Q x = (Q + Q^T)x = 2Qx$  (if  $Q$  is symmetric)
  - $\nabla_x x^T v = \nabla_x v^T x = v$
- Minima or Maxima or ...



- Minimum, maximum, saddle point: things become quickly complicated in high dimensions.
- Formally, you show Hessian is positive definite:  $\nabla_x (\nabla_x f)^T$




# CRASH COURSE ON OPTIMIZATION

---

- $f(x; \theta)$  is a strictly convex function of  $x$ , if:

$$\frac{f(x_1; \theta) + f(x_2; \theta)}{2} < f\left(\frac{x_1 + x_2}{2}; \theta\right), \quad \forall x_1, x_2$$



- Then  $f$  has only one local extremum. It is a local minimum, and this is the global minimum.

# CRASH COURSE ON OPTIMIZATION

---

- Back to our setting:

$$f(x; Q, b, c) = x^T Q x - 2b^T x + c$$


- $Q$  is a symmetric positive-definite matrix.
- Multi-variable Quadratic form.
- This is convex. Single extremum which is a minimum.
- Consider eigen-decomposition of  $Q = V\Lambda V^T$ .
  - Columns of  $V$  are eigen-vectors.  $V$  is unitary.
  - $\Lambda$  is diagonal, with eigen-values. All eigenvalues positive.
- $Q = V\Lambda V^T$ ,  $x^T Q x = (Vx)^T \Lambda (Vx) = \sum_i \lambda_i (Vx)_i^2$
- Sum of quadratic terms with all coefficients ( $\lambda_i$ ) positive

# CRASH COURSE ON OPTIMIZATION

---

- Back to our setting:

$$f(x; Q, b, c) = x^T Q x - 2b^T x + c$$



Positive "semi" definite (Eigenvalues are non-negative)

# CRASH COURSE ON OPTIMIZATION

---

- Back to our setting:

$$f(x; Q, b, c) = x^T Q x - 2b^T x + c$$

- Assume  $Q$  is positive definite:

$$\nabla_x f = 0 \rightarrow 2Qx - 2b = 0 \rightarrow Qx = b$$

- $x = Q^{-1}b$

# CRASH COURSE ON OPTIMIZATION

---

## General note on computing $Q^{-1}b$

- Never compute  $Q^{-1}$ , and then multiply by  $b$ .
  - Numerically unstable, more expensive.
- Call `scipy.linalg.solve`:
  - Cholesky / LDL Decomposition:  $Q = L D L^T$
  - Always exists for a positive definite matrix.  $L$  is lower triangular.
  - Solve  $Qx = b \rightarrow LDL^T x = b \rightarrow Ly = b, L^T x = D^{-1}y$

$$\begin{bmatrix} a & 0 & 0 & 0 & \dots \\ q & c & 0 & 0 & \dots \\ d & e & f & 0 & \dots \\ & \vdots & & & \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ \vdots \end{bmatrix}$$

# DENOISING

---

$$X = \arg \min_X \frac{1}{2\sigma^2} \|Y - X\|^2 + R(X)$$

$$X = \arg \min_X X^T Q X - 2b^T X + c$$

- $R(X) = \lambda \sum_n (x[n] - 0.5)^2 = \lambda \|X - 0.5\|^2$
- $Q = \frac{1}{2\sigma^2} I + \lambda I$
- $Q$  is therefore diagonal.
- $Q^{-1}$  involves inverting elements along diagonal.
- Simple to compute  $Q^{-1}b$ .
  - Independent operation on each pixel / element of  $b$ .

# DENOISING

---

$$X = \arg \min_X \frac{1}{2\sigma^2} \|Y - X\|^2 + R(X)$$

$$X = \arg \min_X X^T Q X - 2b^T X + c$$

- $R(X) = \lambda \sum_n [|(G_x * x)[n]|^2 + |(G_y * x)[n]|^2]$
- $R(X) = \lambda(\|A_{gx}X\|^2 + \|A_{gy}X\|^2)$
- Using  $\|Y\|^2 = Y^T Y$ ,  $(AB)^T = B^T A^T$ :
  - $Q = \frac{1}{2\sigma^2} I + \lambda(A_{gx}^T A_{gx} + A_{gy}^T A_{gy})$
  - $b = \frac{1}{2\sigma^2} Y$
- $Q$  is HUGE and not diagonal.
- Can't even form  $Q$ , let alone call `scipy.linalg.solve`
- You could form 'sparse matrix', but we'll get to that later.

# DENOISING

---

- Need to find  $X = Q^{-1}b$  where
  - $Q = \frac{1}{2\sigma^2}I + \lambda(A_{gx}^T A_{gx} + A_{gy}^T A_{gy})$
  - $b = \frac{1}{2\sigma^2}Y$
- Can we diagonalize  $Q$ ?
- YES ! Use the Fourier Transform / Fourier basis  $S$ 
  - $A_{gx} = SD_{gx}S^*$
  - $A_{gx}^T A_{gx} = S|D_{gx}|^2 S^*$
  - $A_{gy}^T A_{gy} = S|D_{gy}|^2 S^*$
  - $I = SS^* = SIS^*$

$|D_g|^2$  denotes  $D_g^* D_g$ .



# DENOISING

---

- Need to find  $X = Q^{-1}b$  where

- $Q = \frac{1}{2\sigma^2}I + \lambda(A_{gx}^T A_{gx} + A_{gy}^T A_{gy})$

- $b = \frac{1}{2\sigma^2}Y$

$$Q = S \underbrace{\left[ \frac{1}{2\sigma^2}I + \lambda(|D_{gx}|^2 + |D_{gy}|^2) \right]}_{\text{Diagonal}} S^*$$

$$QX = b \rightarrow S^* X = \left[ \frac{1}{2\sigma^2}I + \lambda(|D_{gx}|^2 + |D_{gy}|^2) \right]^{-1} S^* b$$

$$F_X[u, v] = \left[ \frac{1}{2\sigma^2} + \lambda(|F_{gx}[u, v]|^2 + |F_{gy}[u, v]|^2) \right]^{-1} \frac{F_Y[u, v]}{2\sigma^2}$$

$$F_X[u, v] = \frac{F_Y[u, v]}{1 + 2\sigma^2\lambda(|F_{gx}[u, v]|^2 + |F_{gy}[u, v]|^2)}$$

- Caveat: Assumes circular convolution

# DE-BLURRING

---

$$X = \arg \min_X \frac{1}{2\sigma^2} \|Y - A_k X\|^2 + \lambda (\|A_{gx} X\|^2 + \|A_{gy} X\|^2)$$

$$X = \arg \min_X X^T Q X - 2b^T X + c$$

- $b = \frac{1}{2\sigma^2} A_k^T Y$
- $Q = \frac{1}{2\sigma^2} A_k^T A_k + \lambda (A_{gx}^T A_{gx} + A_{gy}^T A_{gy})$
- Still diagonalizable by the Fourier Basis

$$Q = S \underbrace{\left[ \frac{1}{2\sigma^2} |D_k|^2 + \lambda (|D_{gx}|^2 + |D_{gy}|^2) \right]}_{\text{Diagonal}} S^*$$

$$QX = b \rightarrow S^* X = \left[ \frac{1}{2\sigma^2} |D_k|^2 + \lambda (|D_{gx}|^2 + |D_{gy}|^2) \right]^{-1} S^* b$$

- $S^* A_k^T Y = S^* (S D_K S^*)^* Y = D_K^* S^* Y$

# DE-BLURRING

---

$$X = \arg \min_X \frac{1}{2\sigma^2} \|Y - A_k X\|^2 + \lambda (\|A_{gx} X\|^2 + \|A_{gy} X\|^2)$$

$$X = \arg \min_X X^T Q X - 2b^T X + c$$

$$F_X[u, v] = \frac{\bar{F}_k[u, v] F_Y[u, v]}{|F_k[u, v]|^2 + 2\sigma^2 \lambda (|F_{gx}[u, v]|^2 + |F_{gy}[u, v]|^2)}$$

- When  $\lambda = 0$ ,  $F_X = F_Y/F_k$ .
- But this is unstable since  $F_k[u, v]$  can be 0 for some  $[u, v]$ .
- We can see that the regularization term in the denominator dominates for  $u, v$  where  $|F_k[u, v]|^2$  is low.
- This is called Wiener filtering.
- Again remember, assumes circular convolution.

# GENERIC RESTORATION

---

$$X = \arg \min_X \sum_n w[n] \|Y[n] - (X * k)[n]\|^2 + R(x)$$

$$X = \arg \min_X \|D_{\sqrt{w}}(Y - A_k X)\|^2 + R(x)$$

$$X = \arg \min_X X^T (A_k^T D_w A_k) X - 2A_k^T D_w Y + R(x)$$

$$X = \arg \min_X X^T Q X - 2b^T X + c$$

- Now,  $Q$  is no longer diagonalized by the Fourier Basis
- No other choice but Cholesky ?
- $Q$  is hard to form, but we can compute  $Q v$  for any  $v$  very easily.  
 $Q v = A_k^T D_w A_k v + \lambda (A_{gx}^T A_{gx} v + A_{gy}^T A_{gy} v)$ 
  - This takes an "image" shaped vector and returns an image shaped vector.
  - Multiplication by  $A_k, A_{gx}, A_{gy}$  is convolution by corresponding kernels.
  - Multiply by  $D_w$  is a point-wise operation.
  - Multiply by  $A_k^T$  is convolution with flipped kernel.

# CONJUGATE GRADIENT

---

- Generic algorithm for solving  $Qx = b$  for symmetric positive definite  $Q$ .
- Useful when you can multiply by  $Q$  but not 'form' it.

## Basic Idea

- For a given set of vectors  $\{p_1, p_2, \dots, p_N\}$ 
  - that are same size as  $x$
  - linearly independent
  - $N = \text{dimensionality of } x$
- We can write any  $x = \sum_i \alpha_i p_i$
- If we also choose the vectors to be 'conjugate' such that  $p_i^T Q p_j = 0$  for  $i \neq j$ :

$$Qx = b \rightarrow p_k^T Qx = p_k^T b \rightarrow \alpha_i p_k^T Q p_k = p_k^T b \rightarrow \alpha_i = \frac{p_k^T b}{p_k^T Q p_k}$$

# CONJUGATE GRADIENT

---

## Iterative Algorithm

- Begin with some guess  $x_0$  for  $x$  (say all zeros)
- $k = 0, r_0 \leftarrow b - Qx_0, p_0 \leftarrow r_0$
- Repeat
  - $\alpha_k \leftarrow \frac{r_k^T r_k}{p_k^T Q p_k}$
  - $x_{k+1} = x_k + \alpha_k p_k$
  - $r_{k+1} = r_k - \alpha_k Q p_k$
  - $\beta_k = \frac{r_{k+1}^T r_{k+1}}{r_k^T r_k}$
  - $p_{k+1} = r_{k+1} + \beta_k p_k$
  - $k = k + 1$

Stop at some measure of convergence. Pre-conditioned variants. Additional reading:

[https://en.wikipedia.org/wiki/Conjugate\\_gradient\\_method](https://en.wikipedia.org/wiki/Conjugate_gradient_method)

# DE-BLURRING

---

What if we did not have a squared regularizer on gradients ?

$$X = \arg \min_X \sum_n \|Y[n] - (X * k)[n]\|^2 + \lambda \sum_n (\|(G_x * X)[n]\| + \|(G_y * X)[n]\|)$$

No longer a quadratic form. ( $\|\cdot\|$  implies absolute value)

**Variable splitting (Divide and Concur) Approach**

$$X = \arg \min_X \min_{\{c_x[n], c_y[n]\}} \sum_n \|Y[n] - (X * k)[n]\|^2 + \lambda \sum_n (\|c_x[n]\| + \|c_y[n]\|) \\ + \beta \left[ \sum_n ((G_x * X)[n] - c_x[n])^2 + ((G_y * X)[n] - c_y[n])^2 \right]$$

Equivalent when  $\beta \rightarrow \infty$

# DE-BLURRING

---

$$X = \arg \min_X \min_{\{c_x[n], c_y[n]\}} \sum_n \|Y[n] - (X * k)[n]\|^2 + \lambda \sum_n (\|c_x[n]\| + \|c_y[n]\|) \\ + \beta \left[ \sum_n ((G_x * X)[n] - c_x[n])^2 + ((G_y * X)[n] - c_y[n])^2 \right]$$

## Iterative Approach

- Begin with some estimate of  $X$ , and a small value of  $\beta$
- Alternate between
  - Minimizing wrt  $c_x, c_y$  keeping  $X$  constant. Pointwise.
  - Minimizing wrt  $X$  keeping  $c_x, c_y$  constant. Quadratic / Fourier diagonalized.
  - While increasing the value of  $\beta$

Further Reading: [Krishnan and Fergus. Fast Image Deconvolution using Hyper-Laplacian Priors, NIPS 2009.](#)  
Also see the ADMM algorithm.



# COLOR

---



# COLOR

---

Remember, at each pixel:

$$X_r[n] = \int_{\lambda} L(\lambda, n) \Pi_r(\lambda) d\lambda$$
$$X_g[n] = \int_{\lambda} L(\lambda, n) \Pi_g(\lambda) d\lambda$$
$$X_b[n] = \int_{\lambda} L(\lambda, n) \Pi_b(\lambda) d\lambda$$

- $L(\lambda, n)$  is the light incident at  $n$ 
  - We've folded in spatial sensitivity, quantum efficiency, ignored noise.
- Here  $\Pi_r, \Pi_g, \Pi_b$  are the wavelength-dependent transmissions of the camera's color filters.
  - Often called color matching functions.
- Assume these are RAW images (no post-processing).

# COLOR

---

Remember, at each pixel:

$$X_r[n] = \int_{\lambda} L(\lambda, n) \Pi_r(\lambda) d\lambda$$
$$X_g[n] = \int_{\lambda} L(\lambda, n) \Pi_g(\lambda) d\lambda$$
$$X_b[n] = \int_{\lambda} L(\lambda, n) \Pi_b(\lambda) d\lambda$$

## Observations

- This is "projection" of a continuous valued function to three numbers.
  - Loss of information.
  - Metamerism:  $L(\lambda)$  that have the same RGB values.

# COLOR

---

Remember, at each pixel:

$$\begin{aligned}X_r[n] &= \int_{\lambda} L(\lambda, n) \Pi_r(\lambda) d\lambda \\X_g[n] &= \int_{\lambda} L(\lambda, n) \Pi_g(\lambda) d\lambda \\X_b[n] &= \int_{\lambda} L(\lambda, n) \Pi_b(\lambda) d\lambda\end{aligned}$$

## Observations

- Rationale: Models the human visual system.
  - We only have three kind of photoreceptors
  - The standard R,G,B filters "span" the same subspace as human observers.
    - Determined using psycho-physical experiments
    - By the International Commission on Illumination (CIE) in 1931
    - Introduced the concept of primary colors
    - Defined the CIE standard observer

We can't distinguish between metamers either.

# COLOR

---

Remember, at each pixel:

$$X_r[n] = \int_{\lambda} L(\lambda, n) \Pi_r(\lambda) d\lambda$$
$$X_g[n] = \int_{\lambda} L(\lambda, n) \Pi_g(\lambda) d\lambda$$
$$X_b[n] = \int_{\lambda} L(\lambda, n) \Pi_b(\lambda) d\lambda$$

## Observations

- $L(\lambda)$  is the spectrum of the light that reaches the camera.
  - This is a function of both the object surface, and the illumination
  - Lights can be of different colors
  - But human perception of color is very stable under changing illumination
  - **"Color Constancy"**
- Also means metamerism is illumination dependent  
Two objects could have identical RGB values under one light but not another.