

CSE 559A: Computer Vision



[credit: danjodon.deviantart.com]

Fall 2017: T-R: 11:30-1pm @ Lopata 101

Instructor: Ayan Chakrabarti (ayan@wustl.edu).

Staff: Abby Stylianou (abby@wustl.edu), Jarett Gross (jarett@wustl.edu)

<http://www.cse.wustl.edu/~ayan/courses/cse559a/>

Oct 24, 2017

GENERAL

- Problem Set 3 Due Thursday.
- Project Proposals Due Sunday.
- Regular Office Hours this week.

GLOBAL OPTIMIZATION (RECAP)

$$d = \arg \min_d \sum_n C[n, d[n]] + \lambda \sum_{(n, n') \in E} S(d[n], d[n'])$$

- Discrete optimization of **disparity map** d , each $d[n] \in \{0, 1, \dots, D - 1\}$
- $C[n, d[n]]$ comes from our matching cost. How well $L[x, y]$ matches $R[x - d[x, y], y]$.
- E is the set of all pairs of "neighboring" pixel locations.
- S is a function that indicates a preference for $d[n]$ and $d[n']$ to be the same.
- Example 1: 0 if $d[n'] = d[n]$, 1 otherwise.
- Example 2: $|d[n'] - d[n]|$
- Example 3:
 - 0 if $d[n'] = d[n]$
 - T_1 if $|d[n'] - d[n]| < \epsilon$
 - T_2 otherwise.

GLOBAL OPTIMIZATION

$$d = \arg \min_d \sum_n C[n, d[n]] + \lambda \sum_{(n,n') \in E} S(d[n], d[n'])$$

Iterated Conditional Modes

- Begin with $d_0 = \arg \min_d C[n, d[n]]$
- At each iteration t , compute d_{t+1} from d_t , by solving for each pixel in d_{t+1} assuming neighbors have values from d_t .

$$d_{t+1}[n] = \arg \min_{d_n} C[n, d_n] + \lambda \sum_{(n,n') \in E_n} S(d_n, d_t[n'])$$

Does it converge ?

- No Guarantee.
A modified version would converge to a local minima if in each iteration, we only updated one pixel.

GLOBAL OPTIMIZATION

$$d = \arg \min_d \sum_n C[n, d[n]] + \lambda \sum_{(n,n') \in E} S(d[n], d[n'])$$

Iterated Conditional Modes (slow!)

- Begin with $d_0 = \arg \min_d C[n, d[n]]$
- At each iteration t , compute d_{t+1} from d_t , by solving for **one** pixel in d_{t+1} assuming neighbors have values from d_t .

$$d_{t+1}[n_{t+1}] = \arg \min_{d_n} C[n_{t+1}, d_n] + \lambda \sum_{(n_{t+1}, n') \in E_{n_{t+1}}} S(d_n, d_t[n'])$$

Does it converge ?

- No Guarantee.
A modified version would converge to a local minima if in each iteration, we only updated one pixel n_t at iteration t .

GLOBAL OPTIMIZATION


$$d = \arg \min_d \sum_n C[n, d[n]] + \lambda \sum_{(n,n') \in \mathbf{E}} S(d[n], d[n'])$$

- These kind of cost functions / optimization problems are quite common in vision.
- The cost can be interpreted as a log probability distribution:

$$p(d) \propto \prod_n \exp(-C[n, d[n]]) \prod_{(n,n') \in \mathbf{E}} \exp(-\lambda S(d[n], d[n']))$$

- Joint distribution over all the $d[n]$ values.


GLOBAL OPTIMIZATION



$$p(d) \propto \prod_n \exp(-C[n, d[n]]) \prod_{(n, n') \in \mathbf{E}} \exp(-\lambda S(d[n], d[n']))$$

- Joint distribution over all the $d[n]$ values.

GLOBAL OPTIMIZATION




$$p(d) \propto \prod_n \exp(-C[n, d[n]]) \prod_{(n, n') \in \mathbf{E}} \exp(-\lambda S(d[n], d[n']))$$

Question: Are $d[n]$ and $d[n']$ independent if:

- If $(n, n') \in \mathbf{E}$ -- pixels are neighbors?

Reminder: Two variables are independent if we can express their joint distribution as a product of distributions on each variable.

GLOBAL OPTIMIZATION




$$p(d) \propto \prod_n \exp(-C[n, d[n]]) \prod_{(n, n') \in \mathbf{E}} \exp(-\lambda S(d[n], d[n']))$$

Question: Are $d[n]$ and $d[n']$ independent if:

- If $(n, n') \in \mathbf{E}$ -- pixels are neighbors. No
- If $(n, n') \notin \mathbf{E}$ -- pixels are not neighbors ?

GLOBAL OPTIMIZATION




$$p(d) \propto \prod_n \exp(-C[n, d[n]]) \prod_{(n, n') \in \mathbf{E}} \exp(-\lambda S(d[n], d[n']))$$

Question: Are $d[n]$ and $d[n']$ independent if:

- If $(n, n') \notin \mathbf{E}$ -- pixels are not neighbors ? NO. Unless n, n' are parts of disconnected components of graph.

GLOBAL OPTIMIZATION



$$p(d) \propto \prod_n \exp(-C[n, d[n]]) \prod_{(n, n') \in \mathbf{E}} \exp(-\lambda S(d[n], d[n']))$$


Question: Are $d[n]$ and $d[n']$ independent if:

- If $(n, n') \notin \mathbf{E}$, "conditioned" on all the neighbors of n being observed. $p(d[n], d[n'] | \{d[n'']\})$

YES. This is the Markov property. And these kinds of graphical models are called Markov random fields.

Graph structure encodes "conditional independence".

GLOBAL OPTIMIZATION




$$d = \arg \max_d p(d) = \arg \min_d \sum_n C[n, d[n]] + \lambda \sum_{(n,n') \in \mathbf{E}} S(d[n], d[n'])$$

GLOBAL OPTIMIZATION

$$d = \arg \min_d \sum_n C[n, d[n]] + \lambda \sum_{(n,n') \in E} S(d[n], d[n'])$$

- Iterated Conditional Modes really slow.
- No guaranteed solution for arbitrary graphs.
- But could solve it if our graph were a chain (or more generally a tree).


GLOBAL OPTIMIZATION



$$\sum_x C[x, d[x]] + \lambda \sum_x S(d[x], d[x + 1])$$

The total cost of those blocks and the edges was the least.

GLOBAL OPTIMIZATION



Say we only had two nodes:


$$d_1, d_2 = \arg \min_{d_1 d_2} C[1, d_1] + C[2, d_2] + \lambda S(d_1, d_2)$$

$$d_2 = \arg \min_{d_2} \min_{d_1} C[1, d_1] + C[2, d_2] + \lambda S(d_1, d_2)$$

This is the d_2 corresponding to the optimal path

$$\sum_x C[x, d[x]] + \lambda \sum_x S(d[x], d[x + 1])$$

GLOBAL OPTIMIZATION



Say we only had two nodes:


$$d_1, d_2 = \arg \min_{d_1 d_2} C[1, d_1] + C[2, d_2] + \lambda S(d_1, d_2)$$

$$d_2 = \arg \min_{d_2} C[2, d_2] + \min_{d_1} (C[1, d_1] + \lambda S(d_1, d_2))$$

This is a function of d_2
or a table of values for each possible
value of d_2

$$\sum_x C[x, d[x]] + \lambda \sum_x S(d[x], d[x + 1])$$

GLOBAL OPTIMIZATION



$$d_1, d_2, d_3 = \arg \min C[1, d_1] + C[2, d_2] + C[3, d_3] + \lambda S(d_1, d_2) + \lambda S(d_2, d_3)$$


$$d_3 = \arg \min_{d_3} C[3, d_3] + \min_{d_2} \left[\lambda S(d_2, d_3) + C[2, d_2] + \min_{d_1} [\lambda S(d_1, d_2) + C[1, d_1]] \right]$$

This is precisely what we computed for the 2 node case

Also note that once you have this, you don't care about what the value of d_1 was in the inner minimization.

$$\sum_x C[x, d[x]] + \lambda \sum_x S(d[x], d[x + 1])$$

GLOBAL OPTIMIZATION



$$d_1, d_2, d_3 = \arg \min C[1, d_1] + C[2, d_2] + C[3, d_3] + \lambda S(d_1, d_2) + \lambda S(d_2, d_3)$$


$$d_3 = \arg \min_{d_3} C[3, d_3] + \min_{d_2} \left[\lambda S(d_2, d_3) + C[2, d_2] + \min_{d_1} [\lambda S(d_1, d_2) + C[1, d_1]] \right]$$

$$\bar{C}[2, \cdot]$$

$$\bar{C}[3, \cdot]$$

$$\bar{C}[x + 1, d] = C[x + 1, d] + \min_{d'} \lambda S(d, d') + \bar{C}[x, d']$$

GLOBAL OPTIMIZATION



We go from left to right, and doing an arg min on the last \bar{C} gives us the disparity of the last node.


And then we backtrack to find the full chain.

Store best d' for each d .

$$z[x + 1, d] = \arg \min_{d'} \lambda S(d, d') + \bar{C}[x, d']$$

$$\bar{C}[x + 1, d] = C[x + 1, d] + \min_{d'} \lambda S(d, d') + \bar{C}[x, d']$$

GLOBAL OPTIMIZATION



Forward

$$\bar{C}[0, d] = C[0, d]$$

$$z[x + 1, d] = \arg \min_{d'} \lambda S(d, d') + \bar{C}[x, d']$$

$$\bar{C}[x + 1, d] = C[x + 1, d] + \min_{d'} \lambda S(d, d') + \bar{C}[x, d']$$

Backward

$$d[x_{end}] = \arg \min_d \bar{C}[x_{end}, d]$$

$$d[x] = z[x + 1, d[x + 1]]$$

GLOBAL OPTIMIZATION

We could apply this on individual epipolar lines.



GLOBAL OPTIMIZATION

- That's why we want to use a full 2D grid.
- But forward-backward only works on chains (or graphs without cycles).

One flavor of approximate algorithms apply the same idea of forming a $\bar{C}[x, d]$

- TRW-S
- Loopy Belief Propagation
- SGM

GLOBAL OPTIMIZATION

Semi-Global Matching

$$\bar{C}[x, d] = C[x, d] + \min_{d'} \bar{C}[x - 1, d'] + \lambda S(d, d')$$

This is going left to right in the horizontal direction.

Idea: Compute different \bar{C} along different directions ...

and average.

GLOBAL OPTIMIZATION

Semi-Global Matching

$$\bar{C}_{lr}[n, d] = C[n, d] + \min_{d'} \bar{C}_{lr}[n - [1, 0]^T, d'] + \lambda S(d, d')$$

$$\bar{C}_{rl}[n, d] = C[n, d] + \min_{d'} \bar{C}_{rl}[n + [1, 0]^T, d'] + \lambda S(d, d')$$

$$\bar{C}_{du}[n, d] = C[n, d] + \min_{d'} \bar{C}_{du}[n - [0, 1]^T, d'] + \lambda S(d, d')$$

$$\bar{C}_{ud}[n, d] = C[n, d] + \min_{d'} \bar{C}_{ud}[n + [0, 1]^T, d'] + \lambda S(d, d')$$

$$d[n] = \arg \min_d \bar{C}_{lr}[n, d] + \bar{C}_{rl}[n, d] + \bar{C}_{ud}[n, d] + \bar{C}_{du}[n, d]$$

GLOBAL OPTIMIZATION

Semi-Global Matching

