

CSE 559A: Computer Vision



[credit: danjodon.deviantart.com]

Fall 2017: T-R: 11:30-1pm @ Lopata 101

Instructor: Ayan Chakrabarti (ayan@wustl.edu).

Staff: Abby Stylianou (abby@wustl.edu), Jarett Gross (jarett@wustl.edu)

<http://www.cse.wustl.edu/~ayan/courses/cse559a/>

Oct 31, 2017

GENERAL

- Project proposals due today. Submit on Blackboard.
- PSET 4 will be posted today. Due two weeks from now.

OPTICAL FLOW




Slides/ Examples from Steve Seitz, Lana Lazebnik, Subhransu Maji, Yasu Furukawa, The Sintel Dataset from MPI

OPTICAL FLOW

Optical Flow = How did the pixels move in the image plane between two frames ?


$$I_t[x, y] \rightarrow I_{t+1} [x + u[x, y], y + v[x, y]]$$



OPTICAL FLOW

Useful for:

Analyzing motion / shape




OPTICAL FLOW

Useful for:

Analyzing motion / shape

Tracking Objects / People across time



OPTICAL FLOW

Useful for:

Analyzing motion / shape

Tracking Objects / People across time

Image Morphing

$$I_t[x, y] \rightarrow I_{t+1} [x + u[x, y], y + v[x, y]]$$

$$I_t[x + u[x, y]/2, y + v[x, y]/2]$$

Forms a mid-point of the deformation between two ~~frames~~ ...
images


OPTICAL FLOW

Useful for:

Analyzing motion / shape

Tracking Objects / People across time

Image Morphing



OPTICAL FLOW

Useful for:

Analyzing motion / shape

Tracking Objects / People across time

Image Morphing



OPTICAL FLOW

$$I_t[x, y] \rightarrow I_{t+1}[x + u[x, y], y + v[x, y]]$$



Correspondence search, without the benefit of epipolar geometry.

OPTICAL FLOW

Let's try to solve it assuming $u[x, y]$, $v[x, y]$ are very small. (Very little movement).

Also assume "brightness constancy":

$$I[x, y, t] = I[x + u[x, y], y + v[x, y], t + 1]$$

Do a Taylor approximation to "linearize" the RHS.

$$I[x, y, t] = I[x + u, y + v, t + 1] \approx I[x, y, t + 1] + \frac{\partial}{\partial x} I[x, y, t]u + \frac{\partial}{\partial y} I[x, y, t]v$$

$$\frac{\partial}{\partial t} I[x, y, t] + \frac{\partial}{\partial x} I[x, y, t]u + \frac{\partial}{\partial y} I[x, y, t]v \approx 0$$

$$I_t + \langle [I_x, I_y], [u, v] \rangle = 0$$

Lucas-Kanade Method

OPTICAL FLOW

Lucas-Kanade Method

$$I_t + \langle [I_x, I_y], [u, v] \rangle = 0$$


- I_t is an image-shaped array of time gradients (subtracting $I[x, y, t + 1] - I[x, y, t]$)
 - (Not to be confused with I_t notation from the last slide)
- I_x and I_y are x- and y- gradient images (e.g., use Sobel filters).
 - Often, you want to apply these on $\frac{I[x,y,t+1]+I[x,y,t]}{2}$

This is an equation on $u[x, y], v[x, y]$ at each pixel location.

But one equation for two variables. Doesn't tell us about flow vector in direction "orthogonal" to image gradient.


OPTICAL FLOW

Aperture Problem



OPTICAL FLOW

Aperture Problem



OPTICAL FLOW

Aperture Problem



http://en.wikipedia.org/wiki/Barberpole_illusion

OPTICAL FLOW

Lucas-Kanade Method

Solution: Assume u, v is constant in a region, and get multiple equations.

So for $u[x, y] = u, v[x, y] = u$, consider a bunch of x', y' in a window around x, y .

$$I_x[x', y'] u + I_y[x', y'] v = -I_t[x', y']$$

Multiple equations, two variables: solve in the least squares sense.

$$\begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

OPTICAL FLOW

Lucas-Kanade Method

$$\begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

Summations are in a window around x, y .

How would you do this without looping over pixels ?

- Compute $I_x^2, I_x I_y, I_y^2, I_t I_x, \dots$ point-wise.
- Use convolutions to do local summation.
- Form each element of left matrix and right vector as a separate image.
- Invert by pointwise operations on these images.

OPTICAL FLOW

Lucas-Kanade Method

$$\begin{bmatrix} \sum I_x^2 + \epsilon & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 + \epsilon \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

Summations are in a window around x, y .

When will you get good answers and when will you get bad answers ?

In other words, when is the matrix invertible ?

- Matrix will be all zero in a smooth region.
- Matrix will be rank 1 if all gradients in one direction.
- Good when you have general texture.

For stability, add a small value to the diagonal elements of the matrix.

OPTICAL FLOW

Lucas-Kanade Method

Pyramid / Hierarchical Variant to handle large displacements

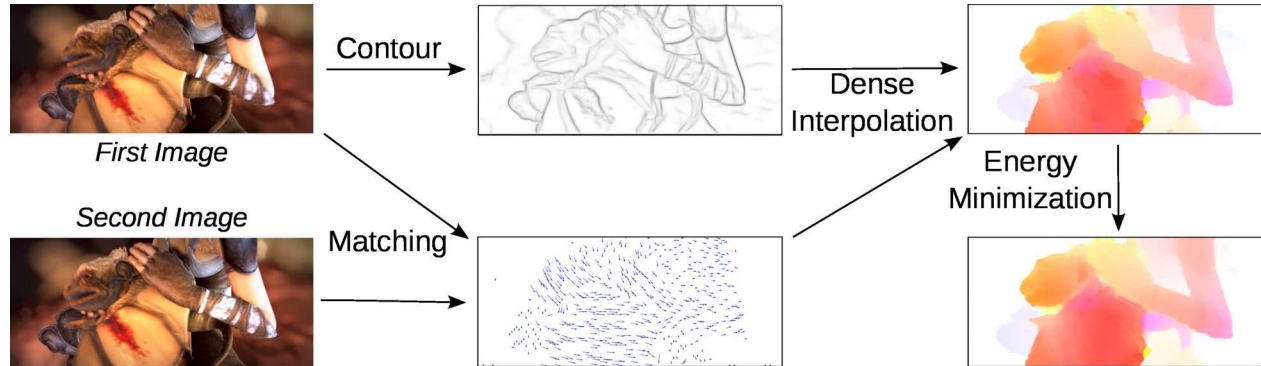
- First downsample images, solve it at a coarser scale.
- Then, upsample the flow-field, and compute displacement from that flow field.

So basically, warp image I_{t+1} based on flow-field from coarser level, and now find differential motion beyond that.

OPTICAL FLOW

State-of-the Art Methods

- Use energy minimization, complex features, contours, ...



Jerome Revaud, Philippe Weinzaepfel, Zaid Harchaoui and Cordelia Schmid
EpicFlow: Edge-Preserving Interpolation of Correspondences for Optical Flow
CVPR 2015.


<http://sintel.is.tue.mpg.de/results> - For a leaderboard of best results.

GROUPING & SEGMENTATION

Partition the set of pixels into disjoint sets or groups



Dual of the edge detection problem !



GROUPING & SEGMENTATION

But what is the basis of this grouping ?

- Physical
 - Lie on the same surface / plane
 - Made of the same material
 - Moving together rigidly
- Semantic
 - Same object
 - Foreground / background
 - Interesting / non-interesting

Semantic segmentation: often humans will disagree on what goes where.

GROUPING & SEGMENTATION

Simplest Version: Superpixel Segmentation

- Partition Image into a large number of segments called superpixels.
- Many segments, each segment relatively small.
- Oversegmentation of the image
 - Each object / plane / surface might be broken into multiple segments
 - But (hope) each segment does not cross a boundary.
- Can be based on appearance alone
- Simplifies further processing (dealing with K segments instead of N pixels)

GROUPING & SEGMENTATION



Image



Superpixels
(different levels)

GROUPING & SEGMENTATION

SLIC Superpixels

Achanta et al., 2010. Simple Linear Iterative Clustering.

Formally, given an image $I[n]$ with N pixels, you want group the pixels into $K \ll N$ super pixels.

You want to determine a label $L[n] \in \{1, 2, \dots, K\}$ for every pixel n , based on some metric.

Note the value of L doesn't matter. What matters is similar pixels have the same label. This is clustering !

GROUPING & SEGMENTATION

SLIC Superpixels

Define an "augmented" image $I'[n]$ where each $I'[n] \in \mathbb{R}^5$

- First 3 dimensions are R,G,B
- Two dimensions are x and y co-ordinates.

For grayscale images, $I'[n] \in \mathbb{R}^3$.

GROUPING & SEGMENTATION

SLIC Superpixels

Determine labeling $L[n]$ to minimize the following cost:

$$L = \arg \min_L \min_{\{\mu_k\}} \sum_{k=1}^K \sum_{n:L[n]=k} \|I'[n] - \mu_k\|^2$$

Here, each $\mu_k \in \mathbb{R}^5$.

- This is K-means clustering.
- Easy to see that μ_k will be the mean of the I' vectors of pixels assigned to label k .
- We're saying that all pixels assigned the label k should be close to each other in the squared distance sense of their augmented vectors.
- This augmented vector encodes both appearance and location.
- So we want pixels that look the same and are close-by to have the same label.

GROUPING & SEGMENTATION

SLIC Superpixels

$$L = \arg \min_L \min_{\{\mu_k\}} \sum_{k=1}^K \sum_{n:L[n]=k} \|I'[n] - \mu_k\|^2$$

- Typically, use Lab color space instead of RGB.
- You can weight the contribution of location vs appearance by normalizing (x, y) in I' differently.

$$I'[n] = [I[n]_R, I[n]_G, I[n]_B, \alpha n_x, \alpha n_y]^T$$

GROUPING & SEGMENTATION

$$L = \arg \min_L \min_{\{\mu_k\}} \sum_{k=1}^K \sum_{n:L[n]=k} \|I'[n] - \mu_k\|^2$$

K-Means: Lloyd's algorithm

- Begin with some initial assignment $L[n]$ (more later).
- At each iteration ...

Step 1: For each k , assign

$$\mu_k = \text{Mean}\{I'[n]\}_{L[n]=k}$$

Step 2: For each n , assign

$$L[n] = \arg \min_k \|I'[n] - \mu_k\|^2$$

- Does this converge ?

But kind of expensive ? In our case, K and N will both be large.

GROUPING & SEGMENTATION

Next time:

- How to initialize $L[n]$, restrict search space of minimization.
- Other kinds of segmentation.